


N-body problem

- N-body simulations are a class of computational problems where we calculate the effect of a force between N objects (or bodies)
- The problem is to calculate the positions and movements of a number of bodies in space as time advances
 - the bodies are affected by all other bodies via gravitation
 - long-range interactions: each object affects all other objects
- Here we will only consider the two-dimensional case
 - extension to three dimensions is straight forward
- We will only consider gravitational force
 - attraction between two bodies
 - can also use the same methods for other types of forces, for instance electrostatic attraction or repulsion

Problem description

- We have N bodies in 2-dimensional space
 - bodies are treated as point masses
 - shape and size does not affect behaviour
- Each body is described by its
 - mass m
 - position $X = (x_x, x_y)$
 - velocity $V = (v_x, v_y)$
the rate of change in position over time, $V = dX/dt$
 - acceleration $A = (a_x, a_y)$
the rate of change in velocity over time, $A = dV/dt$
- As a body is affected by a force, its velocity changes
 - Newton's laws describe how bodies in space affect each other with gravitation



force ← - - - - - →
velocity → - - - - - →

Laws of gravitation


■ Newton's laws

- two bodies a and b with masses m_a and m_b
- positions are $X_a = (x_a, y_a)$ and $X_b = (x_b, y_b)$
- gravitational force on a caused by b

- $F_{a,b} = \frac{Gm_a m_b}{r^2} \frac{X_b - X_a}{r}$ where G is the gravitational constant, $G = 6.67259\text{e-}11$

- and $r = \sqrt{(x_b - x_a)^2 + (y_b - y_a)^2}$ is the distance between a and b


- forces are vectors with x- and y-components $F = (f_x, f_y)$
- pairwise forces are symmetric but of opposite direction: $F_{a,b} = -F_{b,a}$



Laws of gravitation (cont.)

- The total force on a body a is the sum of the pairwise forces from all the other bodies

$$- \quad F_a = \sum_{j=0, j \neq a}^{N-1} F_{a,j} = G m_a \sum_{j=0, j \neq a}^{N-1} m_j \left[\frac{X_j - X_a}{r_{a,j}^3} \right]$$



- The force on a body affects its motion according to Newton's second law
 $F = m * a$
 - $a = F/m$ (acceleration is force divided by mass)
- Given the current position of a body (x_a, y_a) and the acceleration we can compute the velocity and position of the body in the next time step

Discrete solution

- We divide the time into short time intervals of length h (often denoted Δt)
 - starting from an initial state at time t_0 , we calculate the position and velocity for each body at times t_1, t_2, t_3, \dots
 - if the current timestep is t_i then the next timestep is $t_{i+1} = t_i + h$
- $O(N^2)$ algorithm
 - for each time interval
 - for each body b
 - for all other bodies c
 - calculate the force on b caused by c
 - calculate new velocity for b
 - calculate new position for b
- The time interval h has to be short enough to give an accurate solution
 - the amount of computation increases with a shorter interval

Discrete formulation

- For a body with mass m we compute the force F affecting it at time t
 - sum of the forces from all other bodies
- Then we compute the velocity and position of the body for the next timestep, i.e. at time $t+h$
 - new velocity is $V^{t+1} = V^t + A^t h$ where $A^t = \frac{F}{m}$
 - new position is $X^{t+1} = X^t + V^t h$
- When the bodies move to new positions, the forces change and the computation has to be repeated
- Called Euler's method

Sequential solution

- A straightforward sequential solution with Euler's method


```
for (t=0; t<maxt; t++) {           /* For each timestep */
    for (i=0; i<nbodies; i++) {     /* For each body */
        F = Force(i);               /* Force on i */
        V[i] = V_old[i]+F*deltat/m; /* New velocity */
        X[i] = X_old[i]+V_old[i]*deltat; /* New position */
    }

    copy V to V_old;
    copy X to X_old;
}
```

- F , V and X are vectors in the 2-dimensional space
 - ◆ vector multiplication and addition
- We keep the values from previous time step in the arrays V_old and X_old


Problem with Euler's method

- In Euler's method, we assume that the acceleration is constant during the whole time interval
 - we use the acceleration at time t to calculate the new velocity at time $t+h$
- Acceleration is not constant during a time interval
 - when a body gets closer to another body, its acceleration increases
 - our approximation is based on the velocity at time t
- Have to use a small time step to get reliable results




The Leapfrog method

- In the Leapfrog method we use the *mid-point* of the interval to calculate the approximation of the velocity during the interval
 - this is used to calculate the positions at the next point in time
- Velocities and positions are not updated at the same point of time
 - update positions at the *beginning* of each time step and velocities at the *middle* of the time steps
- Positions are updated at times $h, 2h, 3h, \dots$
 - $X_{i+1} = X_i + h \cdot V_{i+1/2}$ for $i=0, 1, 2, \dots$
- Velocities are updated at times $1/2h, t+1/2h, 2t+1/2h, 3+1/2h, \dots$
 - $V_{i+1/2} = V(t+1/2h)$



Properties of the Leapfrog method

- Using the velocity at the midpoint of the interval gives a better approximation than the velocity at the beginning (or end) of the interval
- No more complicated to implement than the Euler scheme
- The initial velocity at $V_{1/2}$ can for instance be calculated by Eulers method
- The leapfrog method is time reversible
 - starting from any state at time t_i we can calculate backwards in time to t_0
- The leapfrog method is a second order approximation
 - the approximations of the positions have an accuracy of $O(\Delta t^2)$



Implementation of the Leapfrog method

- A sequential implementation of the Leapfrog method is

```
set initial values of V_old to velocities at time 0.5*deltat


for (t=0; t<maxt; t++) {           /* For each timestep */
    for (i=0; i<nbodies; i++) {     /* For each body */
        X[i] = X_old[i]+V_old[i]*deltat; /* New position */
        F = Force(i);                /* Force on i */
        V[i] = V_old[i]+F*deltat/m;    /* New velocity */
    }

    copy V to V_old;
    copy X to X_old;
}
```

- We now first update positions, then acceleration and velocities

Parallel N -body simulation: domain decomposition


- N bodies, P processes
 - we divide the bodies evenly among the processes
 - $M = N/P$ local bodies per process



- A process computes the force that affects its *local* bodies
 - for this, it needs information about all the other bodies
- When a process has updated the values for its own bodies, the results are communicated to all processes
 - the updated positions, velocities and acceleration are needed in the next time step

Communication structure

- There are different ways to implement the distributed computation and communication
- One possible solution is to connect the processes in a ring
 - circulate the local bodies of each process through the ring
 - after $P-1$ steps, all bodies have visited all processes and all force interactions have been calculated
- In a system with point-to-point connected processors, communication can proceed simultaneously between all processes
 - all processes have the same amount of work, both computation and communication
- Another alternative solution is to do the exchange of data with collective communication



Algorithm structure

■ Algorithm for a ring-based solution

- for each time interval
 - update positions of local bodies
 - calculate pairwise force between local bodies
 - for all other processes in the ring
 - send local bodies to next process
 - receive bodies from previous process
 - compute pairwise force between the received and the local bodies
 - update velocities of local bodies

■ The force-values are accumulate with contributions of particles from the other processes

- all these are added together to get the total forces

Ring algorithm for process k

- Each process has neighbours *next* and *previous*
- Local bodies are stored in an array B , non-local bodies are received in an array *inbuf*
- *ComputeForce* computes forces between two sets of bodies

```
next      = (k==P-1) ? 0 : k+1;  /* Next and previous */
previous  = (k==0) ? P-1 : k-1;  /* process in the ring */

for (t=0; t<maxt; t++) {        /* For each time step */
    update positions of local particles
    ComputeForce(B, B);          /* Forces between local bodies */
    for (step=1; step<P; step++) { /* Loop through ring */
        send B to next;
        receive inbuf from previous; /* Forces between */
        ComputeForce(B, inbuf);      /* nonlocal bodies */
    }
    update velocities for particles
}
```

Implementation with collective communication

- We can also use collective communication for the communication
- Use MPI collective communication to exchange information about the bodies to all processes
 - the communication does not need to do the data exchange in $P-1$ steps, but can directly send the needed data to all processes
 - the implementation is otherwise similar, but the communication structure becomes more simple

